

Concourse: The Event Loop

Concourse research notes

July 17, 2026

Contents

1	What this document is	1
2	D1: fire emits clock deltas	2
2.1	Why the diff loses	2
2.2	The delta vocabulary	2
2.3	The debug oracle	2
2.4	Refinement forced by F2: deltas are derived, not emitted	2
3	D2 proposal: the cascade worklist	3
3.1	The problem	3
3.2	Split contention from propagation	3
3.3	What falsifies it	4
4	D3: the interpreter owns the record	4
5	D4: the builder chooses the sampler	5
6	The interpreter, whole	5
7	Charter extensions	5

1 What this document is

`model_definition.tex` §8 queued four middle-layer questions. Design discussion on July 17, 2026 settled three and framed the fourth:

- D1 (decided).** `fire` emits enable/disable/re-enable deltas; the full `enabled(state)` recompute is retained as a debug-mode consistency check, not measured against — deltas are chosen for correctness of the order of operations, not for speed.
- D2 (open; proposal in §3).** The cascade worklist must settle to a fixed point; the policy that makes the fixed point unique is proposed here and awaits a decision.
- D3 (decided).** The interpreter owns the marked record. `CompetingClocks` samples, and only samples; its watchers observe sampling, and the marked record is not a sampling artifact.

D4 (decided). Sampler choice belongs to CompetingClocks’ `SamplerBuilder`: Concourse describes the clock-key type and the distributions in play, and the builder picks. Any sampler valid for those distributions must work; none is load-bearing.

Charter claims F9–F11 (§7) extend the falsification charter of `model_definition.tex` §7.

2 D1: fire emits clock deltas

2.1 Why the diff loses

Beyond ordering, an enabled-set diff is *semantically blind to re-enablings*. When processor sharing re-evaluates every resident clock on an arrival, or a preemption changes which distribution a still-enabled key carries, the enabled set before and after the firing is identical — the diff sees nothing, while the sampler must receive `reenable!`. Deltas carry information the diff cannot express. The diff survives only as the debug oracle for the one thing it can check: membership.

2.2 The delta vocabulary

The contract’s `fire` is θ -free and pure, so deltas cannot carry distributions — only keys and operations. The interpreter derives each distribution from `clock_distribution` at the post-firing state and each anchor from `st.te` (amendment A1):

```
const ClockDelta = Tuple{Symbol,ClockKey}    # (:enable / :disable / :reenable,
      key)
# fire returns (new_state, deltas::Vector{ClockDelta}); emission order is
# authoritative and is applied to the sampler verbatim.
```

This is ClockGradients’ existing `fire_changes(model, state, key)` seam with the changes type made concrete (and the CD-1 time argument applied to it as well); `enabled_update` consumes the same deltas, so the estimators’ incremental path and the interpreter share one vocabulary. The fired clock itself is implicit: the interpreter calls `fire!(ctx, key, t)` before applying deltas, so `fire` never emits a delta for its own key.

2.3 The debug oracle

After every firing, debug mode asserts two things:

1. *Membership*: `enabled(m, st')` equals the previous enabled set plus emitted `:enables` minus emitted `:disables` (and the fired key). This is the complete-recompute check, and it is the specification — `enabled` remains the meaning, deltas the implementation.
2. *Anchors*: the sampler’s per-clock enabling times equal `st.te` — the A1 invariant. This is what checks re-enablings, which membership cannot see.

2.4 Refinement forced by F2: deltas are derived, not emitted

Running charter claim F2 (add reneging touching only the family registry and the surface constructor) falsified D1 *as first implemented*: patience clocks are born when a job enters a buffer and die when it leaves, and with deltas hand-emitted inside `settle!` and the family

fire functions, adding the family would have meant threading emission code through every state-transition site — interpreter surgery, which is exactly what F2 forbids. The repair keeps D1’s decision and changes its mechanism:

- *Membership deltas are derived, not emitted.* The cascade already knows which stations it touched. After the cascade, for each touched station and each family, the family’s per-station enabled keys are diffed between the old and new states (both exist; **fire** is pure). A family is now exactly four methods — **family_station_keys!**, **family_low**, **family_fire!**, **family_memory** — and clock lifecycle follows from membership alone.
- *Locality is preserved.* The user’s original argument for deltas over a global diff was cost on large networks; the derived diff runs only over the touched stations’ clocks, so the asymptotics stand. Explicit **:reenable** deltas remain available for the processor-sharing path (phase 2), which membership cannot see.
- *The A1 bookkeeping collapses into one place.* Enabling times and banked ages are updated in the same derivation pass, driven by the diffs and the per-family memory policy, instead of at every site that moves a job.
- *A correctness dividend.* A clock enabled and disabled within one firing (a job dispatched and immediately preempted in the same cascade) now nets out of the diff and never reaches the sampler — the GSMP has no zero-duration clocks, which the emission order previously only approximated by a disable/enable pair.

This is the falsification charter working as designed: the claim failed against the first implementation, and what survived is a better mechanism under the same decision. After the repair, the **:patience** addition itself changed only the registry and the surface fields, and $M/M/1+M$ matches the Erlang-A birth–death form with the membership oracle running through every renegeing cascade.

3 D2 proposal: the cascade worklist

3.1 The problem

Within one firing at time t , instantaneous consequences chain: a deposit triggers **settle!** at the destination; a departure from a full buffer lets a blocked upstream server transfer in, which frees that server, which admits from its buffer, which may unblock stations further upstream. The cascade must reach a fixed point — no station can dispatch, preempt, or unblock — before the sampler sees any delta. The requirement is that the fixed point be *unique and deterministic*, because **fire** is the replayed, differentiated object.

The GSPN literature has this exact problem as immediate transitions and vanishing markings, resolved there by priorities and weights (Zimmermann Ch. 5–6). The proposal below is that resolution, restated for queues.

3.2 Split contention from propagation

The order of cascade processing matters in exactly one place: when two consumers compete for one freed resource — two blocked upstream servers and one freed buffer slot; two idle servers and

one arriving job is already settled by the discipline’s `select`. Everywhere else the steps merely propagate.

Contention is model semantics. Who gets a freed slot is an observable modeling choice, not an implementation detail, so it must be a declared policy value on the station, parallel to the discipline:

```

UnblockPolicy:
  FCFSUnblock()           # longest-blocked first (default)
  PriorityUnblock(by)      # by a ScalarExpr over the blocked job's marks
  RandomUnblock()         # a recorded choice, via the A4 draw source

```

`FCFSUnblock` needs the blocking order, so the blocked-server set in `QueueState` is an ordered vector, not a set. `RandomUnblock` consumes a recorded draw and so stays inside the A4 randomness census.

Propagation must be order-free. With contention delegated to policy, the claim is that the remaining cascade relation is *confluent*: any fair processing order reaches the same fixed point. Termination comes from check C3 (acyclic blocking chains): each unblocking transfer moves a job strictly downstream in the blocking DAG, so the cascade measure is well-founded. Confluence plus termination give a unique normal form (Newman’s lemma); the worklist order is then an implementation choice, not a semantic one.

The canonical order, so that the implementation is deterministic even if the confluence claim fails: a FIFO worklist of stations to settle, seeded with the destination station then the origin station of the firing — the one-hop order `queue_layers.tex` §3.2 fixed by fiat, generalized. A debug-build fuel counter bounds the cascade at (jobs + servers) steps and turns nontermination into an error rather than a hang.

3.3 What falsifies it

F9 in the charter: run the same trajectories with FIFO and LIFO worklists; `states_equal` at every index or the confluence claim is dead, and the worklist order gets promoted into the IR’s semantics the way the one-hop order already is. Either outcome is a result; the test is cheap.

4 D3: the interpreter owns the record

The marked record (k_i, t_i, a_i) is the primary observable of the whole library, and its auxiliary-draw column exists because the *interpreter* supplies the draw source (A4) — the sampler never sees those draws, so a sampler-side watcher could never capture them. The interpreter appends one entry per firing: the key, the time, and the draws consumed by that firing in consumption order. CompetingClocks’ watchers remain what they are — instrumentation of sampling — and Concourse attaches none of the record path to them. `gradient_record` gets one method, from the marked record type, per `model_definition.tex` §6.

One replay detail follows from record ownership: during estimation, `fire` must read firing i ’s recorded draws, and the contract does not pass i . `QueueState` therefore carries a firing counter (`nfired::Int`, incremented by every `fire`), and the replay binding indexes the record with it. A pure counter in a pure state; the same device as `next_id`.

5 D4: the builder chooses the sampler

Concourse tells `SamplerBuilder` the clock-key type and lets `CompetingClocks` choose the sampler for the distributions in play; the default is whatever the builder decides, and Concourse hard-codes no method. The design consequence is a claim, not a configuration: *every* sampler valid for the model’s distribution set must produce statistically identical results, because the interpreter speaks only the blessed `SamplingContext` surface (`enable!`, `disable!`, `reenable!`, `next`, `fire!`) and no model code reaches around it. Estimator-driven overrides (the carry coupling for IPA) are passed through as builder arguments by the estimation entry points, not chosen by the model.

6 The interpreter, whole

With D1–D4 fixed, the middle layer is small enough to state completely. Julia-flavored pseudocode; error handling and the debug oracle elided:

```
function simulate(m::QueueGSMP,  $\theta$ , horizon, seed)
  st = initial_state(m)
  ctx = SamplingContext(build(m), seed)      # D4: builder chooses
  rec = MarkedRecord(clockkeytype(m))
  for k in enabled(m, st)                   # initial enablings
    enable!(ctx, k, clock_distribution(m,  $\theta$ , k, st), st.te[k], 0.0)
  end
  while true
    (t, key) = next(ctx, st.time)
    (key === nothing || t > horizon) && break
    fire!(ctx, key, t)                      # realize the winner's draw
    draws = draw_source(ctx, key)           # A4: keyed, recording
    st, deltas = fire_changes(m, st, key, t, draws) # cascade inside
    push!(rec, key, t, consumed(draws))      # D3: interpreter-owned
    for (op, k) in deltas                   # D1: emission order verbatim
      op === :disable && disable!(ctx, k, t)
      op === :enable && enable!(ctx, k, clock_distribution(m,  $\theta$ , k, st),
        st.te[k], t)
    end
  end
  close_record!(rec, min(horizon, st.time))
end
```

Everything model-specific is inside `fire_changes` (families, disciplines, kernels, the cascade); everything stochastic is inside `ctx` and `draws`; everything observable is `rec`. The loop itself should never grow.

7 Charter extensions

Continuing the numbering of `model_definition.tex` §7:

F9 (cascade confluence). With contention resolved by `UnblockPolicy`, the cascade fixed point is order-independent. *Test:* FIFO vs. LIFO worklists over random finite-buffer `:block`

networks; `states_equal` at every firing index. Falsified $\Rightarrow$ the worklist order becomes IR semantics.

F10 (sampler agnosticism). No charter result depends on the sampler. *Test:* run the F1–F7 oracle comparisons under at least two builder outcomes (e.g. first-to-fire and next-reaction); all pass under both.

F11 (delta fidelity). Emitted deltas agree with the recompute oracle. *Test:* debug-mode membership and anchor assertions (§2) hold over long mixed runs including preemption — and, in phase 2, processor sharing, which is the case only the anchor check can see.